


Sparsity in NN: Sources and Structures

Diverse Sparsity Ranges/Patterns Must be Exploited Well

- Computer Vision Models:** ~40%, 70%-80% in later layers of deep CNNs
 - Weight sparsity: 50% (MobileNetV2) – 93% (EfficientNets)
 - 80% - 85% in Point-wise Convolutions
 - Can be structured (dimension/block pruning, fixed density in a block)
 - Activation sparsity is low (~20% and can be) - unstructured
- Language Models:** ~80% – 93% (baseline Transformers, BERTs, etc.)
 - But, existing accelerators do not exploit 80%+ sparsity well.
 - Very coarse-grain in large language models (e.g., Switch transformers)
 - Easy to exploit; Pruning unimportant tokens and heads (Up to 75%)
- RNNs:** Up to ~60% activation sparsity; ~80% weight sparsity
 - Usually unstructured, but can be structured with pruning or with factorized operators (batch norm, quantization, activation function)
- Drop-out Layers**
- Atrous (dilated convolutions):** fine-grained structured in Weights
- GANs:** ~60% in Activations, in transposed CONV in degenerators.
- 3D Point Clouds:** Up to 80% or more in Activations, Unstructured.
- Gradient Sparsity in Communication:** >> 95% Unstructured
 - More than 99% for computer vision or language tasks
 - 95% – 99% for recommendation models
 - Challenging to exploit – both storage-wise and compute-wise
- Graph Learning:** High (75%–99%) or Hyper (99%+) Unstructured
 - Local features lead to sparse adjacency matrices and sparsity propagates (sparse/dense, dense/dense multiplications)
 - Dense/Sparse compute processed with separate accelerator modules
- Text Analysis:** ~67% – 99% Unstructured; Weights (and Activations)




- Accelerating sparsity is important for many other domains**
 - linear algebra, graph processing, scientific computing, database, genomics, compression (Usually unstructured)

Need Special Hardware or Software Mechanisms

- Sparsity cannot be leveraged as it is, including on DNN accelerators**
- Need special mechanisms (Even for structured sparsity) for
 - Encode:** Store only non-zeros with their locations (Encode)
 - Decode:** Get non-zero values from on/off-chip memory or storage.
 - Extract:** Find matching non-zeros from two tensors to multiply/add
 - Load Balance:** Ensure each computing unit has similar work
 - Communicate:** Both non-zero and its position
 - Do All of Above **Without Much Overhead** (Power, Performance, Area)

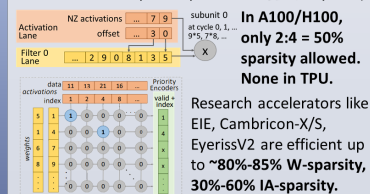
Effectiveness of Special Mechanisms

Works Great for Limited Range in ~30%-80%





- Storage Benefits With Structured Sparsity:**
 - Same as unstructured, for density-bounded blocks (2:4 in A100/H100)
 - Benefits closer to ideal as coarse-grain pruning is higher (e.g., dimensions)

- Non-Zero Matching for scalar/vectors:**
 - Lookup:** Use position of a non-zero from first tensor to extract non-zero from second tensor
 - Intersection:** Compare streams of positions of non-zeros from two tensors
 - Size of a buffer used for lookup/comparison determines sparsity that can be exploited**
 - 16 non-zeros in 256 positions to feed 16-input vector unit exploits up to 1-1/16 ~93.75% sparsity

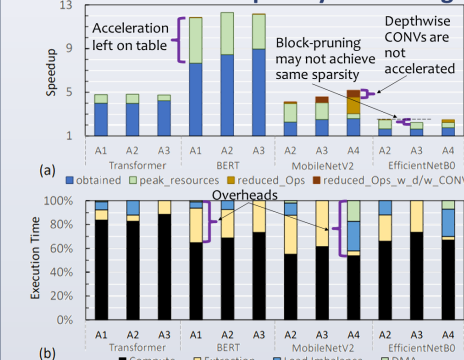


- Structured Sparsity Helps in Load Balancing:**
 - Without Balancing, Work among Leading and Trailing Compute Units Differ Noticeably** (~6.5X)
 - Without balancing the work, idle cycles for computational units can be 40%-50%.



- Data Reuse (arithmetic intensity) Reduces With Sparsity**
 - So are opportunities for less on-chip communication (spatial reuse), low memory accesses (temporal reuse)
 - Reuse in CONV2D and batched GEMMs/MLPs still considerable

Overheads of Sparsity Processing



- Energy reductions relatively quite easier to attain** – lowering memory accesses is the dominating factor.
- Latency reductions are efficient up to only ~80% sparsity**
- Bottlenecks:** Non-zero matching or imbalance—even when weights pruned in structure in tandem with unstructured activations
 - ~8X speedup out of 12X achievable for 92% sparse BERT
 - ~50% of achievable speedup for compact vision models (MobileNets)
 - ~4X lower speedup attained vs. achievable at hyper sparsity
 - Imbalance makes infeasible to feed large vector/SIMD units
 - Hard to achieve similar sparsity with block-pruning (EfficientNets)
- Most accelerators poorly process depthwise CONV – slowdowns**


Designing All-Sparsity Exploiting Accelerator

- Reconfigurable Mechanisms Needed for Flexibility**
 - Sparsity-tailored multiple encodings and non-zero matchings
 - Communication interconnects for multicast, unicast of data and spatial/temporal accumulations
 - Unified Memory for data and metadata
 - Configurable workgroup formulation and asynchronous processing
- Novel mechanisms for high & hyper sparsity (>90%, especially 99%+)**

Algorithmic Sparsity Can Be Made System-Aware

- Apply quantization and sparsity (pruning, operator reformation) based on **hardware's capability**
- Leverage execution models of the system** – find out what advantage a sparsity or quantization can offer on your target hardware/compiler
 - A model with 80% uniform sparsity across layers may perform better than a model containing layers with 70% and 90% sparsity, with same accuracy.
- Maximize improvements jointly with quantization, pruning, and value similarity (interplay on compression, acceleration exploited, and accuracy)
- Cross-layer Codesign

Quantitative Approach for Cross-layer Design




AutoML for Model Sparsity and Compression

- For expert-directed or automated search of best compression or codesign, **specify all common hyperparameters for applicability**
 - A variety of pruning options (unstructured, 1D or k:n block-sparsity, bit-widths of tensors/layer, tolerable accuracy, goals: storage/energy/performance)
- Automated optimization for hyperparameters for generality and efficiency**
 - Pruning ratio for each iteration (epoch); pruning mechanism (which values to prune, e.g., below a certain threshold); pruning pattern (fine-grain, block size); bit-widths of tensors (quantization).
 - This does not need to be manual or explored from a limited pre-designed set.

Tensor Core Accelerators Exploit Sparsity Better

- Sparsity below 99% usually do not lead to practical speedups on CPUs/GPUs without tensor cores
 - Recent algorithmic & software advances strive to reduce gap
- Tensor-cores-based accelerators can provide much higher performance/watt (TPUs, A100)
- Trends and Roadmap:** Flexible hardware/compiler and Cross-layer Codesign for Sparsity



[Gale, Zaharia, Young, Elsen, SC'20]